

```

#include <pthread.h>
#include <iostream>
#include <map>
#include <vector>
#include <string.h>
#include <signal.h>
#include <set>
using namespace std;

//!
//! \class critical_section_speculator
//!
//! \brief implements a class that allows multiple threads to process concurrently shared,
//! data, while keeping a sequential consistency related to the order in which this
//! threads invoke the speculation. This order is user's responsibility.
//!
//! Through this class instead of using locks, each thread can go speculatively into
//! it's critical section, assuming that it will consume the values produced by the
//! previous threads affecting the data and none of the values produced by later threads.
//!
//! In order to do this the class relies on write_data and read_data functions,
//! that check for data dependence violations. As well as a per-thread required
//! commit function, with which the thread signals that it has ended it's use of
//! the critical section.
//!
//! Limitations:
//! * If a speculative thread is to be canceled, it cannot use functions
//! that involve system mutexes, such as printf, etc. In this case, it
//! is possible that the thread can be canceled while holding such a mutex,
//! and the application can go into deadlock. In order to prevent this the
//! user has to surround this "dangerous" code with:
//! "pthread_setcancelstate(PTHREAD_CANCEL_DISABLE, NULL);" and
//! "pthread_setcancelstate(PTHREAD_CANCEL_ENABLE, NULL);".
//!
//! * Sequential consistency is mostly guaranteed, save for exception
//! behaviour.
//!
class critical_section_speculator {
private:

    //bool values to set if the speculator is active or if the shared_data
    //has already been initialized.
    bool _is_active, _has_fixed_shared_data;

    //mutex for synchronized use of the previous variable
    pthread_mutex_t _is_active_mutex;

    //the speculative threads and it's data
    vector<_loop_spec_thread> _spec_threads;

    //red-black tree, used as a thread index to relate a pthread_id with
    //it's model id, i.e. it's position in _spec_threads
    map<pthread_t, int> _thread_index;

    //mutex for a synchronized access to both of the former
    pthread_mutex_t _spec_threads_mutex;

    //shared_data data between the threads
    void*& _shared_data;

    //auxiliary data to reset _shared_data.
    int _null_data;

    //red-black tree that relates the reference of a data element from
    //shared_data data with a _data_access_log used to keep track of all the
    //possible data dependence violations
    map<void*, _data_access_log> _access_log;

    //mutex related to the _access_log as a whole
    pthread_mutex_t _global_access_log_mutex;

    //! private method that allows to reset the object
    //!
    //! TO BE NOTED:
    //! * this method takes all class mutexes. On invocation
    //! _spec_threads_mutex and _global_access_log_mutex should
    //! be on hold and no other class mutex, save those related
    //! to a thread that is not about to be canceled or restarted.
    //!
    void _reset_speculator () {

```

```

        if (!_spec_threads.empty()){
            for (int i=0; i<_spec_threads.size(); i++){
                pthread_mutex_lock(&_spec_threads[i]._thread_mutex);
            }
        }

        pthread_mutex_lock(&_is_active_mutex);

        if (!_spec_threads.empty()){
            for (int i=0; i<_spec_threads.size(); i++){
                pthread_mutex_destroy(&_spec_threads[i]._thread_mutex);
            }
            _spec_threads.clear();
        }

        if (!_thread_index.empty()){

            /*in order to clear the _thread_index, it will be necessary to guarantee that all the threads in deferred
            cancel have finished, and thus will not use the object, affecting its future behaviour*/

            _is_active=false;

            pthread_mutex_unlock(&_is_active_mutex);

            pthread_mutex_unlock(&_spec_threads_mutex);

            pthread_mutex_unlock(&_global_access_log_mutex);
            map <pthread_t, int>::iterator it;
            for (it=_thread_index.begin(); it!=_thread_index.end(); it++){
                if (it->second==1){
                    int a;
                    do {
                        a= pthread_kill(it->first, 0);
                    } while (a==0);
                }
            }

            pthread_mutex_lock(&_global_access_log_mutex);

            pthread_mutex_lock(&_spec_threads_mutex);

            _thread_index.clear();

            if (!_access_log.empty()){
                _access_log.clear();
            }
        }
        else {

            if (!_access_log.empty()){
                _access_log.clear();
            }

            _is_active=false;

            pthread_mutex_unlock(&_is_active_mutex);
        }
    };

    /// private method that allows to reset speculative threads, which has to be called
    /// from a thread whose position is passed as argument. This method is used for resetting
    /// threads that committed some form of data dependence violation.
    ///
    /// TO BE NOTED:
    /// * this method takes all class mutexes. On invocation
    ///   _spec_threads_mutex and _global_access_log_mutex should
    ///   be on hold and no other class mutex, save those related
    ///   to a thread that is not about to be canceled or restarted.
    ///
    ///
    void _reset_spec_threads (int current_thread, vector <int> threads_to_reset){

        int i;
        bool valid_arguments=true;

        if (!threads_to_reset.empty()){

            for (i=0; i<threads_to_reset.size(); i++){

                if (threads_to_reset[i]<0 || threads_to_reset[i]> static_cast<int>(_spec_threads.size())){
                    valid_arguments=false; //one thread has an invalid id.
                    i=threads_to_reset.size();
                }
            }
        }
    }

```

```

    }
    else{
        valid_arguments=false; //no threads to delete
    }
    if (current_thread<0){
        valid_arguments=false;
    }
    if (valid_arguments){

        set<int> threads_currently_held;
        threads_currently_held.insert(current_thread);

        set<int> set_of_threads_to_reset;
        set_of_threads_to_reset.insert(threads_to_reset.begin(), threads_to_reset.end());

        set<int>:: iterator it;
        it=set_of_threads_to_reset.begin();

        while (it!=set_of_threads_to_reset.end()){

            unsigned int previous_size=set_of_threads_to_reset.size();

            if (threads_currently_held.find(*it)==threads_currently_held.end()){

                pthread_mutex_lock(&_spec_threads[*it]._thread_mutex);
                threads_currently_held.insert(*it);

                set<void*>::iterator it2;

                for (it2=_spec_threads[*it]._read_data.begin(); it2!=_spec_threads[*it]._read_data.end(); it2++){

                    _access_log.find(*it2)->second._readers.erase(*it);

                }

                _spec_threads[*it]._read_data.clear();

                for (it2=_spec_threads[*it]._written_data.begin(); it2!=_spec_threads[*it]._written_data.end(); it2++){

                    void* value_to_restore= _access_log.find(*it2)->second._get_previous_value(current_thread)

                    vector<int> writers_to_cancel=_access_log.find(*it2)->second._cancel_higher_writers(current_thread);
                    vector<int> readers_to_cancel=_access_log.find(*it2)->second._cancel_higher_readers(current_thread);

                    if (!writers_to_cancel.empty()){//implied or secondary WAW and WAR violations have occurred
                        memcpy((void*)&const_cast<void*>(*it2), (void*)&value_to_restore, _access_log.find(*it2)->second._get_size());
                    }

                    if (!readers_to_cancel.empty()){//implied or secondary RAW violations have occurred
                        writers_to_cancel.insert(writers_to_cancel.end(), readers_to_cancel.begin(), readers_to_cancel.end());
                    }

                    if (!writers_to_cancel.empty()){
                        for (unsigned int i=0; i<writers_to_cancel.size(); i++){
                            set_of_threads_to_reset.insert(writers_to_cancel[i]);
                        }
                    }

                }

                _spec_threads[*it]._written_data.clear();

                if ((*it!=current_thread)&&(pthread_equal(_spec_threads[*it]._thread, pthread_self())==0)){

                    pthread_mutex_lock(&_is_active_mutex);

                    _spec_threads[*it]._commit=false;

                    if (pthread_kill(_spec_threads[*it]._thread, 0)==0){
                        pthread_cancel(_spec_threads[*it]._thread);
                    }

                    _thread_index[_spec_threads[*it]._thread]=-1;

                    pthread_mutex_unlock(&_is_active_mutex);

                }

            }

            if (set_of_threads_to_reset.size()>previous_size){
                it=set_of_threads_to_reset.begin();
            }
            else {
                it++;
            }
        }

        pthread_mutex_lock(&_is_active_mutex);

```

```

        pthread_attr_t attr;
        pthread_attr_init (&attr);
        pthread_attr_setschedpolicy(&attr, SCHED_RR);

        for (it=set_of_threads_to_reset.begin(); it!=set_of_threads_to_reset.end(); it++){

            if (*it!=current_thread) {

                pthread_create(&_spec_threads[*it]._thread, &attr, _spec_threads[*it]._thread_instructions, _spec_
                _thread_index[_spec_threads[*it]._thread]= *it;

            }

            pthread_mutex_unlock(&_is_active_mutex);

            for (it=threads_currently_held.begin(); it!=threads_currently_held.end(); it++){
                if (*it!=current_thread)
                    pthread_mutex_unlock(&_spec_threads[*it]._thread_mutex);
            }

        }

};

public:

    ///
    /// default constructor
    ///
    critical_section_speculator():_shared_data((void*)&_null_data){

        _is_active=false;
        _null_data=-1;
        _has_fixed_shared_data=false;
        pthread_mutex_init (&_is_active_mutex, NULL);
        pthread_mutex_init (&_spec_threads_mutex, NULL);
        pthread_mutex_init (&_global_access_log_mutex, NULL);

    };

    ///
    /// function providing access to the shared_data as a whole
    ///
    void*& get_shared_data () {

        pthread_mutex_lock(&_is_active_mutex);

        if (!_is_active){

            pthread_mutex_unlock(&_is_active_mutex);

            return (void*)&_null_data; //the object is inactive

        }

        pthread_mutex_unlock(&_is_active_mutex);

        pthread_mutex_lock(&_spec_threads_mutex);

        if (_thread_index.empty() || _thread_index.find(pthread_self())!=_thread_index.end()){

            pthread_mutex_unlock(&_spec_threads_mutex);

            return _shared_data; //valid thread or thread in deferred cancelation

        }

        pthread_mutex_unlock(&_spec_threads_mutex);

        return (void*)&_null_data; // invalid caller

    };

    ///
    /// required function that signals that a given thread has ended
    /// its execution
    ///
    void commit () {

        pthread_mutex_lock(&_is_active_mutex);

        if (_is_active){

            pthread_mutex_unlock(&_is_active_mutex);

            pthread_mutex_lock(&_spec_threads_mutex);

```

```

        if (!_thread_index.empty()){ //thread in deferred cancel.
            pthread_t thread_id=pthread_self();

            if (_thread_index.find(thread_id)!=_thread_index.end()){

                int pos=_thread_index.find(thread_id)->second;
                if (pos>=0){ //valid thread
                    pthread_mutex_lock(&_spec_threads[pos]._thread_mutex);
                    _spec_threads[pos]._commit=true;
                    pthread_mutex_unlock(&_spec_threads[pos]._thread_mutex);
                }
            }

            pthread_mutex_unlock(&_spec_threads_mutex);
        }
        else{
            pthread_mutex_unlock(&_is_active_mutex);
        }
    };

    ///
    /// function to be called from the speculative threads
    /// in order to read the shared_data while keeping the
    /// expected sequential data consistency
    ///
    template <class T>
    T read_data(T*& data_to_be_read){

        pthread_mutex_lock(&_is_active_mutex);

        if (_is_active){

            pthread_mutex_unlock(&_is_active_mutex);

            int current_thread=-2;
            pthread_mutex_lock(&_spec_threads_mutex);

            pthread_t thread_id=pthread_self();

            if (_thread_index.empty()){ //thread in deferred cancel.
                current_thread=-3;
            }
            else if (_thread_index.find(thread_id)!=_thread_index.end()){
                current_thread=_thread_index.find(thread_id)->second;
                if (current_thread===-1) //thread in deferred cancel.
                    current_thread=-3;
            }
            else {
                current_thread=-3; //invalid caller
            }
            if (current_thread>=-1) //is a valid read.

                T retval;

                pthread_mutex_lock(&_global_access_log_mutex);
                if (_access_log.find((void*)data_to_be_read)==_access_log.end()){
                    //if the data has never been accesed, it's log should be created

                    pthread_mutex_t temp;
                    pthread_mutex_init (&temp, NULL);
                    _data_access_log temp2;

                    //since it's the data's first use, it's assumed to be the value in the pre-critical-section
                    //so the data is initialized on that premise
                    temp2._size=sizeof(T);
                    temp2._writers.push_back((int)-1);
                    temp2._previous_values.push_back((void*)*(T*)data_to_be_read);

                    if (current_thread!==-1){
                        pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
                        _spec_threads[current_thread]._read_data.insert((void*)data_to_be_read);
                        temp2._readers.insert(current_thread); //the current read is logged
                    }
                    pthread_mutex_unlock(&_spec_threads_mutex);

                    _access_log.insert (pair<void*, _data_access_log>((void*)data_to_be_read, temp2));
                    retval=*(T*)data_to_be_read;

                    if (current_thread!==-1){
                        pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
                    }
                    pthread_mutex_unlock(&_global_access_log_mutex);

```

```

        return retval;
    }
    _data_access_log* ptr=&_access_log.find((void*)data_to_be_read->second);
    pthread_mutex_lock(&ptr->_data_mutex);
    pthread_mutex_unlock(&_global_access_log_mutex);

    retval= (T) ptr->_get_previous_value(current_thread);
    vector<int> threads_to_cancel= ptr->_cancel_higher_writers(current_thread);

    //if the log exists, then anti-dependence violation is checked (WAR)

    if (current_thread!=-1){
        pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);

        if (!threads_to_cancel.empty()){//an anti-dependence violation has occurred
            memcpy((void*)data_to_be_read, (void*)&retval, sizeof(T)); //the previous value is restored.
            _reset_spec_threads(current_thread, threads_to_cancel);
        }
        else{
            retval=*(T*) data_to_be_read;//no anti-dependence violation has occurred, the data can be read on i
        }
        pthread_mutex_unlock(&_spec_threads_mutex);

        //now that the reading is done, it has to be logged
        if (current_thread!=-1){
            ptr->_readers.insert(current_thread);
            _spec_threads[current_thread]._read_data.insert((void*)data_to_be_read);
            pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
        }
        pthread_mutex_unlock(&ptr->_data_mutex);
        return retval;
    }
    pthread_mutex_unlock(&_spec_threads_mutex);
}
else {
    pthread_mutex_unlock(&_is_active_mutex);
}
return *(T*) data_to_be_read; //invalid caller, inactive object or thread in deferred cancel.
//it's calling data is returned instead of NULL, to prevent segmentation faults.
};

//!
//! function to be called from the speculative threads
//! in order to write the shared_data
//! while keeping the expected sequential data consistency
//!
template <class T>
int write_data(T*& data_to_be_written_upon, T* data_to_write){
    pthread_mutex_lock(&_is_active_mutex);
    if (_is_active){
        pthread_mutex_unlock(&_is_active_mutex);

        int current_thread=-2;

        pthread_mutex_lock(&_spec_threads_mutex);

        pthread_t thread_id=pthread_self();

        if (_thread_index.empty()){//thread in deferred cancel.
            current_thread=-3;
        }
        else if (_thread_index.find(thread_id)!=_thread_index.end()){
            current_thread=_thread_index.find(thread_id)->second;
            if (current_thread==-1)//thread in deferred cancel.
                current_thread=-3;
        }
        else {
            current_thread=-3; //invalid caller
        }
        if (current_thread>=-1){

            pthread_mutex_lock(&_global_access_log_mutex);

            if (_access_log.find((void*)data_to_be_written_upon)==_access_log.end()){
                //if data has no log entry, then it's entry should be created.
                pthread_mutex_t temp;
                pthread_mutex_init (&temp, NULL);
                _data_access_log temp2;

                //since it's the data's first use, it's assumed to be the initial value
                temp2._size=sizeof(T);
                temp2._writers.push_back((int)-1);
            }
        }
    }
}

```

```

temp2._previous_values.push_back((void*)*(T*)data_to_be_written_upon);

if (current_thread!=-1){
    pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
    _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
    temp2._writers.push_back(current_thread);
    temp2._previous_values.push_back((void*)*(T*)data_to_be_written_upon);
}
else{
    temp2._writers.push_back(-1);
    temp2._previous_values.push_back((void*)data_to_write);
}
pthread_mutex_unlock(&_spec_threads_mutex);

_access_log.insert (pair<void*, _data_access_log>((void*)data_to_be_written_upon, temp2));

//and finally the writting is done
memcpy ((void*)data_to_be_written_upon, (void*)&data_to_write, sizeof(T));

if (current_thread!=-1){
    pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
}

pthread_mutex_unlock(&_global_access_log_mutex);
return 0;
}

//if the log exists, then possible true dependence and output dependence violations are checked (RAW and WAW)

_data_access_log* ptr=&_access_log.find((void*)data_to_be_written_upon)->second;
pthread_mutex_lock(&ptr->_data_mutex);
pthread_mutex_unlock(&_global_access_log_mutex);

if (current_thread!=-1){
    pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
}

T value;
vector<int> writers_to_cancel;

bool no_output_violation=false;
if ((void*)*(T*)data_to_be_written_upon==(void*)data_to_write){
    //Since restarting threads has such a significant cost, it is better to check if it is really necessary to
    //So, we will evaluate false positives in output dependence violations.

    if (!ptr->_output_violations_are_false(data_to_write, current_thread)){

        value= (T) ptr->_get_previous_value(current_thread);
        writers_to_cancel=ptr->_cancel_higher_writers(current_thread);

    }
    else {
        no_output_violation=true;
    }
}
else{
    value= (T) ptr->_get_previous_value(current_thread);
    writers_to_cancel=ptr->_cancel_higher_writers(current_thread);
}

vector<int> readers_to_cancel=ptr->_cancel_higher_readers(current_thread);

if (no_output_violation && readers_to_cancel.empty()){
    if (current_thread!=-1){
        _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
        pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
    }
    pthread_mutex_unlock(&ptr->_data_mutex);
    pthread_mutex_unlock(&_spec_threads_mutex);
    return 0;
}

if (!writers_to_cancel.empty()){
    //an output dependence violation has occurred, it's previous value has to be restored before writting over,
    //those writers have to be restarted
    ptr->_writers.push_back(current_thread);
    ptr->_previous_values.push_back((void*)value);
}

/*This branch has been turned into a comment, since the logging for all cases where there are false output
is done on _output_violations_are_false(..). However, should a unforeseeable case arise where this does not
(for instance, a particular thread reset pattern), then it should be un-commented.

else if (!no_output_violation){
    //no output dependence violation, it's current value is valid and logged as the previous value of
    //thread write...
    ptr->_writers.push_back(current_thread);
}

```

```

        ptr->_previous_values.push_back((void*)(T*)data_to_be_written_upon);
    }*/

    if (!readers_to_cancel.empty()){//a true dependence violation has occurred, those readers should be restart
        writers_to_cancel.insert(writers_to_cancel.end(), readers_to_cancel.begin(), readers_to_cancel.end()
    }
    if (!writers_to_cancel.empty()){
        //the violating threads are restarted
        _reset_spec_threads(current_thread, writers_to_cancel);
    }
    pthread_mutex_unlock(&_spec_threads_mutex);
    //now some more logging of the write
    if (current_thread!=-1){
        ptr->_writers.push_back(-1);
        ptr->_previous_values.push_back((void*)data_to_write);
    }
    else {
        _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
    }

    //and finally the write is made
    memcpy ((void*)data_to_be_written_upon, (void*)&data_to_write, sizeof(T));

    if (current_thread!=-1){
        pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
    }
    pthread_mutex_unlock(&ptr->_data_mutex);
    return 0;
}
pthread_mutex_unlock(&_spec_threads_mutex);
else {
    pthread_mutex_unlock(&_is_active_mutex);
}
return -1; //invalid caller, thread in deferred cancel or inactive object.
};

/**
 * speculate: a function that takes the instructions and arguments of a critical section
 * of a given thread and starts it's execution alongside with other threads in their
 * own related critical sections, while maintaining the sequential consistency.
 * This function modifies the global &shared_data, and returns a void*, representing
 * additional per-thread values the user might return from the function given as argument.
 */
/**
 * TO BE NOTED: This function waits in an infinite loop for the threads to signal that they
 * have committed.
 */
void* speculate (void*& shared_data, void* (f)(void*), void* const_args_f){

    pthread_mutex_lock(&_spec_threads_mutex);
    pthread_mutex_lock(&_global_access_log_mutex);

    pthread_mutex_lock(&_is_active_mutex);
    if (!_is_active){
        _is_active=true;
        _has_fixed_shared_data=false;
    }
    pthread_mutex_unlock(&_is_active_mutex);

    if (!_has_fixed_shared_data){
        _shared_data=shared_data;
        _has_fixed_shared_data=true;//if the speculator was resetted, the shared data is re-loaded.
    }

    pthread_attr_t attr;
    pthread_attr_init (&attr);
    pthread_attr_setschedpolicy(&attr, SCHED_RR);

    //the thread is created
    int pos=_spec_threads.size();
    _spec_threads.resize(pos+1);
    pthread_mutex_lock(&_spec_threads[pos]._thread_mutex);
    _spec_threads[pos]._thread_instructions=f;
    _spec_threads[pos]._const_args=const_args_f;
    int success=pthread_create (&_spec_threads[pos]._thread, &attr, f, const_args_f);
    if (success!=0){
        pthread_mutex_unlock(&_spec_threads[pos]._thread_mutex);
        _spec_threads.resize(pos);
        pthread_mutex_unlock(&_global_access_log_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);
        return (void*) _null_data; //the thread could not be created
    }
    _thread_index[_spec_threads[pos]._thread]=pos;
    pthread_mutex_unlock(&_spec_threads[pos]._thread_mutex);

```



```

pthread_mutex_unlock(&_global_access_log_mutex);
pthread_mutex_unlock(&_spec_threads_mutex);

for (int i=0; i<pos; i++){
    do {
        } while (!_spec_threads[i]._commit);
}
bool commit_made=false;
void* retval, *prevval;
bool first_time=true;
do{
    if (!first_time){
        prevval=retval;
    }
    else {
        first_time=false;
    }
    success=pthread_join (_spec_threads[pos]._thread, &retval);
    if (success!=0){
        if (_spec_threads[pos]._commit){
            commit_made=true;
            retval=prevval;
        }
        else{
            success=0;
        }
    }
} while (success==0);

pthread_mutex_lock(&_spec_threads_mutex);
pthread_mutex_lock(&_global_access_log_mutex);
if (pos==(static_cast<int>(_spec_threads.size()-1))){ //if it was the last thread then it can reset the speculator.
    _reset_speculator();
}
pthread_mutex_unlock(&_global_access_log_mutex);
pthread_mutex_unlock(&_spec_threads_mutex);

return retval;
};
};

```