

```

#include <pthread.h>
#include <map>
#include <vector>
#include <string.h>
#include <signal.h>
#include <set>

using namespace std;

// Definition of type script_function, that masks a void* (*)(void*)
typedef void* (*script_function)(void*);
// Definition of type script_vector, a std::vector of type script_function
typedef std::vector<script_function> script_vector;

//
// structure: _loop_spec_thread, used to encapsulate a pthread,
//           which will run a speculative thread for loop speculation,
//           keeping as well some related data.
//
struct _loop_spec_thread{
    pthread_mutex_t _thread_mutex;
    pthread_t _thread;

    set <void*> _written_data; //will be used on thread-cancellation,
    //to undo any changes it has made.

    set <void*> _read_data; //will be used on thread-cancellation,
    //to delete it's logs as reader.

    void* (*_thread_instructions) (void*);
    void* _const_args;

    bool _commit; //helps to determine if the loop iteration
    //that the _loop_spec_thread represents has committed or not.

    _loop_spec_thread(){
        pthread_mutex_init (&_thread_mutex, NULL);
        _commit=false;
    };
};

//
// structure: _data_access_log, used to encapsulate the use history of a shared data,
//           item: it's writers, readers and previous values. It is through this
//           structure that data dependence violation is tracked.
//
struct _data_access_log{
    unsigned int _size; //size of data
    set <int> _readers;
    vector <int> _writers;
    vector <void*> _previous_values; //this will be used to solve anti-dependence
    //and output dependence violations (WAR and WAW)
    //restoring the data to a value that is safe.

    pthread_mutex_t _data_mutex;

    _data_access_log(){
        pthread_mutex_init (&_data_mutex, NULL);
    };

    template <class T>
    bool _output_violations_are_false(T* data_to_write, int pos){ //this function determines if the output dependence
    //violations are false, that is, if the value to write is the same that is already written by later iterations.
    //Note that in this version this check to avoid unnecessary thread-resetting is only possible for WAW violations, since
    //there is book-keeping of the values that are being written, but not of the values that are being read.

        void* prev_value;
        int first_reader=-2;;
        for (int i=0; i<_writers.size(); i++){
            if (_writers[i]==-1){
                prev_value=_previous_values[i];
            }
            else if ( _writers[i]>pos){
                if (data_to_write!=(T*)_previous_values[i]){
                    if (first_reader==-2){
                        if (_previous_values[i]!=prev_value){
                            return false;
                        }
                        first_reader=_writers[i];
                    }
                }
            }
            else if (first_reader==_writers[i]){
                if (_previous_values[i]!=prev_value){

```

```

        return false;
    }
    else{
        return false;
    }
}

if (first_reader==-2){
    _writers.push_back(pos);
    _previous_values.push_back((void*)data_to_write);
    return true;
}
for (int i=0; i<_writers.size(); i++){
    if (_writers[i]==first_reader){
        _previous_values[i]=(void*)data_to_write;
    }
}
_writers.push_back(pos);
_previous_values.push_back(prev_value);
return true;
};

vector<int> _cancel_higher_readers(int pos){
    //if used before a write_data and the return vector is not empty, a true-dependence violation has occurred (RAW).
    //and the higher readers have to be restarted.
    vector<int> to_cancel;
    set<int>::iterator it=_readers.begin();
    while (((!_readers.empty()) && it!=_readers.end())){
        if ((*it)>pos){
            to_cancel.push_back(*it);
            _readers.erase(*it);
            it=_readers.begin();
        }
        else {
            it++;
        }
    }
    if ((!to_cancel.empty())&&(!_writers.empty())){
        for (unsigned int i=0; i<to_cancel.size(); i++){
            unsigned int j=0;
            while(((!_writers.empty())&& j<_writers.size())){
                if (to_cancel[i]==_writers[j]){
                    _writers.erase(_writers.begin()+j); //all the writings of the higher readers on this data
                    _previous_values.erase(_previous_values.begin()+j);
                    j=0;
                }
                else
                    j++;
            }
        }
    }
    return to_cancel;
};

void* _get_previous_value(int pos){
    void* value;
    int currPos=-2;
    //The logic to find the previous value is to get the previous value from the lesser of the higher writers,
    //if that fails, then it is necessary to find the previous value of the latest of the lower writers.

    for (int i=(static_cast<int>(_writers.size()-1); i>-1; i--){ //this order matters because if a thread has written several
        //it's first should be restored

        if (static_cast<int>(_writers[i])>pos){
            if (currPos==-2){
                value=_previous_values[i];
                currPos=_writers[i];
            }
            else if (static_cast<int>(_writers[i])<=currPos){
                value=_previous_values[i];
                currPos=_writers[i];
            }
        }
    }

    if (currPos==-2){
        for (int i=(static_cast<int>(_writers.size()-1); i>-1; i--){
            if (static_cast<int>(_writers[i])<=pos){
                if (currPos==-2){
                    value=_previous_values[i];
                    currPos=_writers[i];
                }
                else if (static_cast<int>(_writers[i])>currPos){
                    value=_previous_values[i];
                }
            }
        }
    }
}

```

```

        currPos=_writers[i];
    }
}

}

}

if (currPos!=-2){
    return (void*)value;
}
return (void*)-1;
};

vector<int> _cancel_higher_writers(int pos){
//if used on a write function and the return vector is not empty, an output dependence violation has occurred (WAW) and the
//value has to be restored to it's state before pos, so it can be written over.

//if used on a read function and the return vector is not empty, an anti-dependence violation has occurred (WAR) and the
//value has to be restored to it's state before pos, so it can be read.
    set<int> to_cancel_set;
    unsigned int i=0;
    while (((!_writers.empty())&&(i<_writers.size())){
        if (_writers[i]>pos){
            to_cancel_set.insert(_writers[i]);
            _writers.erase(_writers.begin()+i);
            _previous_values.erase(_previous_values.begin()+i);
            i=0;
        }
        else
            i++;
    }
    vector<int> to_cancel;
    to_cancel.insert(to_cancel.end(), to_cancel_set.begin(), to_cancel_set.end());
    if ((!to_cancel.empty())&&(!_readers.empty())){
        for (unsigned int i=0; i<to_cancel.size(); i++){
            _readers.erase(to_cancel[i]); //all the readings of the higher writers on this data item have to be unlogg
        }
    }
    return to_cancel;
};

};

//!
//! \class loop_speculator
//!
//! \brief implements a class that takes n ordered iterations from a loop and
//! executes them speculatively in parallel, while keeping the sequential
//! consistency. In order to do this the class relies on write_data and
//! read_data functions, that check for data dependence violations, as well
//! as a commit function that helps to verify that an iteration has ended it's
//! execution.
//!
//! Additional functionality is provided with functions like append and
//! valid_til, that allow the pre-loop section to dynamically add new
//! iterations or subtract iterations to an on-going speculative execution.
//! This allows the model to mimic do-while and do-until behaviour.
//!
//! Limitations:
//! * If a speculative thread is to be canceled, it cannot use functions
//! that involve system mutexes, such as printf, etc. In this case, it
//! is possible that the thread can be canceled while holding such a mutex,
//! and the application can go into deadlock. In order to prevent this the
//! user has to surround this "dangerous" code with:
//!     "pthread_setcancelstate(PTHREAD_CANCEL_DISABLE, NULL);" and
//!     "pthread_setcancelstate(PTHREAD_CANCEL_ENABLE, NULL);".
//!
//! * Sequential consistency is mostly guaranteed, save for exception
//! behaviour.
//!
class loop_speculator {
private:

    //bool values to set if the speculator is active, or has control
    //of the pre-loop section
    bool _is_active, _has_pre_loop;

    //mutex for synchronized use of the previous group of variables
    pthread_mutex_t _is_active_mutex;

    //thread to manage the pre-loop section, boolean value to check if
    //it has ended it's execution & related mutex
    pthread_t _pl;

```

```

bool _pl_commit;
pthread_mutex_t _pl_mutex;

//the speculative threads and it's data
vector<_loop_spec_thread> _spec_threads;

//red-black tree, used as a thread index to relate a pthread_id with
//it's model id, i.e. it's position in _spec_threads
map<pthread_t, int> _thread_index;

//mutex for a synchronized access to both of the former
pthread_mutex_t _spec_threads_mutex;

//shared_data data between the pre-loop section & the speculative threads
void*& _shared_data;

//auxiliary data to reset _shared_data.
int _null_data;

//variable indicating until what iteration the pre-loop section has validated
//the execution (-1 if not used) & related mutex
int _valid_til;
pthread_mutex_t _valid_til_mutex;

//red-black tree that relates the reference of a data element from
//shared_data data with a _data_access_log used to keep track of all the
//possible data dependence violations
map<void*, _data_access_log> _access_log;

//mutex related to the _access_log as a whole
pthread_mutex_t _global_access_log_mutex;

//! private method that allows to cancel speculative threads, whose
//! positions are passed as an argument. It is used if there is a
//! error on starting the speculation, and also for resetting the object.
//!
//! TO BE NOTED:
//! * this method takes all class mutexes. On invocation
//!   _spec_threads_mutex and _global_access_log_mutex should
//!   be on hold and no other class mutex, save those related
//!   to a thread that is not about to be canceled or restarted.
//!
//! * if called with an empty vector, it will reset all logs and
//! wait for the threads in deferred cancel to finish.
//!
void _cancel_spec_threads_and_reset_logs (vector<int> threads_to_delete){

    int i;
    bool valid_arguments=true;

    if (!threads_to_delete.empty()){

        for (i=0; i<threads_to_delete.size(); i++){

            if (threads_to_delete[i]<0 || threads_to_delete[i]> static_cast<int>(_spec_threads.size())){
                valid_arguments=false; //one thread has an invalid id.
                i=threads_to_delete.size();
            }

        }

    }
    else{
        valid_arguments=false; //no threads to delete

        /*This option will be used to initialize the object, resetting it's inner arrays*/

        pthread_mutex_lock(&_is_active_mutex);

        pthread_mutex_lock(&_valid_til_mutex);

        if (_has_pre_loop)

            pthread_mutex_lock(&_pl_mutex);

        if (!_spec_threads.empty()){
            _spec_threads.clear();
        }

        if (!_thread_index.empty()){

            /*in order to clear the _thread_index, it will be necessary to guarantee that all the threads in deferred
            cancel have finished, and thus will not use the object*/

```

```

        if (_has_pre_loop)
            pthread_mutex_unlock(&_pl_mutex);

        pthread_mutex_unlock(&_valid_til_mutex);

        pthread_mutex_unlock(&_is_active_mutex);

        pthread_mutex_unlock(&_spec_threads_mutex);

        pthread_mutex_unlock(&_global_access_log_mutex);

        map<pthread_t, int>::iterator it;
        for (it=_thread_index.begin(); it!=_thread_index.end(); it++){
            if (it->second==-1){
                int a;
                do {
                    a= pthread_kill(it->first, 0);
                } while (a==0);
            }
        }

        pthread_mutex_lock(&_global_access_log_mutex);

        pthread_mutex_lock(&_spec_threads_mutex);

        _thread_index.clear();

        if (!_access_log.empty()){
            _access_log.clear();
        }
    }
    else {

        if (!_access_log.empty()){
            _access_log.clear();
        }

        if (_has_pre_loop)

            pthread_mutex_unlock(&_pl_mutex);

        pthread_mutex_unlock(&_valid_til_mutex);

        pthread_mutex_unlock(&_is_active_mutex);

    }

}

if (valid_arguments){

    for (unsigned int i=threads_to_delete.size()-1; i>=0; i--){
        pthread_mutex_lock(&_spec_threads[threads_to_delete[i]]._thread_mutex);
    }

    pthread_t thrd_id=pthread_self();

    pthread_mutex_lock(&_is_active_mutex);
    pthread_mutex_lock(&_valid_til_mutex);

    if (_has_pre_loop)

        pthread_mutex_lock(&_pl_mutex);

    for (unsigned int i=0; i<threads_to_delete.size(); i++){
        if (pthread_equal(_spec_threads[threads_to_delete[i]]._thread, thrd_id)==0){
            if (pthread_kill(_spec_threads[threads_to_delete[i]]._thread, 0)==0){
                pthread_cancel(_spec_threads[threads_to_delete[i]]._thread);
            }
            _thread_index[_spec_threads[threads_to_delete[i]]._thread]= -1;
            pthread_mutex_destroy(&_spec_threads[threads_to_delete[i]]._thread_mutex);
        }
    }

    if (_has_pre_loop)
        pthread_mutex_unlock(&_pl_mutex);

    pthread_mutex_unlock(&_valid_til_mutex);

    pthread_mutex_unlock(&_is_active_mutex);

}

};

```

```

///! private method that allows to cancel speculative threads higher than the
///! one passed as an argument, it is used in the valid_til function.
///!
///! TO BE NOTED:
///! * this method takes all class mutexes. On invocation
///!   _spec_threads_mutex and _global_access_log_mutex should
///!   be on hold and no other class mutex, save those related
///!   to a thread that is not about to be canceled or restarted.
///!
///!
void _cancel_higher_spec_threads (int top_thread_pos){

    bool valid_arguments=false;

    if (top_thread_pos>=-1 && top_thread_pos< static_cast<int>(_spec_threads.size())){

        set <int> threads_to_cancel;
        set <void*> data_to_check;
        int aux_value=top_thread_pos;

        if (aux_value===-1){
            aux_value=0;
        }
        for (unsigned int i=aux_value; i<_spec_threads.size(); i++){
            threads_to_cancel.insert(i);
        }

        set <int>::iterator it;
        for (it=threads_to_cancel.begin(); it!=threads_to_cancel.end(); it++){
            if (*it!=top_thread_pos){
                pthread_mutex_lock(&_spec_threads[*it]._thread_mutex);
            }
        }

        set<void*>::iterator it2;
        for (it=threads_to_cancel.begin(); it!=threads_to_cancel.end(); it++){

            for (it2=_spec_threads[*it]._read_data.begin(); it2!=_spec_threads[*it]._read_data.end(); it2++){

                if (data_to_check.find(*it2)==data_to_check.end()){

                    data_to_check.insert(*it2);
                }
            }

            if (*it!=top_thread_pos){

                _spec_threads[*it]._read_data.clear();
            }

            for (it2=_spec_threads[*it]._written_data.begin(); it2!=_spec_threads[*it]._written_data.end(); it2++){

                if (data_to_check.find(*it2)==data_to_check.end()){

                    data_to_check.insert(*it2);
                }
            }

            if (*it!=top_thread_pos){

                _spec_threads[*it]._written_data.clear();
            }
        }

        pthread_mutex_lock(&_is_active_mutex);
        pthread_mutex_lock(&_valid_til_mutex);
        if (!_has_pre_loop)
            pthread_mutex_lock(&_pl_mutex);

        for (it=threads_to_cancel.begin(); it!=threads_to_cancel.end(); it++){
            if (*it!=top_thread_pos){
                _spec_threads[*it]._commit=false;
                if (pthread_kill(_spec_threads[*it]._thread, 0)==0){
                    pthread_cancel(_spec_threads[*it]._thread);
                }
                _thread_index[_spec_threads[*it]._thread]= -1;
                pthread_mutex_destroy(&_spec_threads[*it]._thread_mutex);
            }
        }

        for (it2=data_to_check.begin(); it2!=data_to_check.end(); it2++){

```

```

        void* new_val= _access_log.find(*it2)->second.get_previous_value(top_thread_pos);
        vector<int> writers_to_cancel=_access_log.find(*it2)->second._cancel_higher_writers(top_thread_pos);
        vector<int> readers_to_cancel=_access_log.find(*it2)->second._cancel_higher_readers(top_thread_pos);

        if (!writers_to_cancel.empty()){
            memcpy((void*)&const_cast<void*>(*it2), (void*)&new_val, _access_log.find(*it2)->second._size);
        }

    }

    if (top_thread_pos==-1){
        _spec_threads.clear();
    }
    else {
        _spec_threads.resize(top_thread_pos+1);
    }

    if (_has_pre_loop)
        pthread_mutex_unlock(&_pl_mutex);
    pthread_mutex_unlock(&_valid_til_mutex);
    pthread_mutex_unlock(&_is_active_mutex);
}

};

/// private method that allows to reset speculative threads, which has to be called
/// from a thread whose position is passed as argument. This method is used for resetting
/// threads that committed some form of data dependence violation.
///
/// TO BE NOTED:
/// * this method takes all class mutexes. On invocation
///   _spec_threads_mutex and _global_access_log_mutex should
///   be on hold and no other class mutex, save those related
///   to a thread that is not about to be canceled or restarted.
///
///
void _reset_spec_threads (int current_thread, void* data_currently_held, vector<int> threads_to_reset){

    int i;
    bool valid_arguments=true;

    if (!threads_to_reset.empty()){

        for (i=0; i<threads_to_reset.size(); i++){

            if (threads_to_reset[i]<0 || threads_to_reset[i]> static_cast<int>(_spec_threads.size())){
                valid_arguments=false; //one thread has an invalid id.
                i=threads_to_reset.size();
            }

        }

    }
    else{
        valid_arguments=false; //no threads to delete
    }

    if (current_thread<-1){
        valid_arguments=false;
    }

    if (valid_arguments){

        set<int> threads_currently_held;
        threads_currently_held.insert(current_thread);

        set<int> set_of_threads_to_reset;
        set_of_threads_to_reset.insert(threads_to_reset.begin(), threads_to_reset.end());

        set<int>::iterator it;
        it=set_of_threads_to_reset.begin();

        set<void*> data_already_checked;
        data_already_checked.insert(data_currently_held);

        while (it!=set_of_threads_to_reset.end()){

            unsigned int previous_size=set_of_threads_to_reset.size();

            if (threads_currently_held.find(*it)==threads_currently_held.end()){

                pthread_mutex_lock(&_spec_threads[*it]._thread_mutex);
                threads_currently_held.insert(*it);

                set<void*>::iterator it2;

                for (it2=_spec_threads[*it]._read_data.begin(); it2!=_spec_threads[*it]._read_data.end(); it2++){

                    _access_log.find(*it2)->second._readers.erase(*it);
                }

            }

        }

    }

}

```

```

_spec_threads[*it]._read_data.clear();

for (it2=_spec_threads[*it]._written_data.begin(); it2!=_spec_threads[*it]._written_data.end(); it2++)
    if (data_already_checked.find(const_cast<void*>(*it2))==data_already_checked.end()){
        data_already_checked.insert(const_cast<void*>(*it2));
        void* value_to_restore= _access_log.find(*it2)->second._get_previous_value(current_thread);
        vector<int> writers_to_cancel=_access_log.find(*it2)->second._cancel_higher_write;
        vector<int> readers_to_cancel=_access_log.find(*it2)->second._cancel_higher_read;
        if (!writers_to_cancel.empty()){//implied or secondary WAW and WAR violations have occurred
            memcpy((void*)&const_cast<void*>(*it2), (void*)&value_to_restore, _access_log.find(*it2)->second._get_previous_value(current_thread));
        }
        if (!readers_to_cancel.empty()){//implied or secondary RAW violations have occurred
            writers_to_cancel.insert(writers_to_cancel.end(), readers_to_cancel.begin(), readers_to_cancel.end());
        }
        if (!writers_to_cancel.empty()){
            for (unsigned int i=0; i<writers_to_cancel.size(); i++){
                set_of_threads_to_reset.insert(writers_to_cancel[i]);
            }
        }
    }
_spec_threads[*it]._written_data.clear();

if ((*it!=current_thread)&&(pthread_equal(_spec_threads[*it]._thread, pthread_self())==0)){
    pthread_mutex_lock(&_is_active_mutex);
    pthread_mutex_lock(&_valid_til_mutex);
    if (_has_pre_loop)
        pthread_mutex_lock(&_pl_mutex);

    _spec_threads[*it]._commit=false;
    if (pthread_kill(_spec_threads[*it]._thread, 0)==0){
        pthread_cancel(_spec_threads[*it]._thread);
    }
    _thread_index[_spec_threads[*it]._thread]=-1;

    if (_has_pre_loop)
        pthread_mutex_unlock(&_pl_mutex);
    pthread_mutex_unlock(&_valid_til_mutex);
    pthread_mutex_unlock(&_is_active_mutex);
}

}
if (set_of_threads_to_reset.size()>previous_size){
    it=set_of_threads_to_reset.begin();
}
else {
    it++;
}
}

pthread_mutex_lock(&_is_active_mutex);
pthread_mutex_lock(&_valid_til_mutex);
if (_has_pre_loop)
    pthread_mutex_lock(&_pl_mutex);

pthread_attr_t attr;
pthread_attr_init(&attr);
pthread_attr_setschedpolicy(&attr, SCHED_RR);

for (it=set_of_threads_to_reset.begin(); it!=set_of_threads_to_reset.end(); it++){
    if (*it!=current_thread) {
        if (_valid_til!=-1){
            if (*it<=_valid_til){
                pthread_create(&_spec_threads[*it]._thread, &attr, _spec_threads[*it]._thread_instructions,
                    _thread_index[_spec_threads[*it]._thread]=*it);
            }
        }
        else {
            pthread_create(&_spec_threads[*it]._thread, &attr, _spec_threads[*it]._thread_instructions,
                _thread_index[_spec_threads[*it]._thread]=*it);
        }
    }
}

```



```

        }
    }

    if (_has_pre_loop)
        pthread_mutex_unlock(&_pl_mutex);

    pthread_mutex_unlock(&_valid_til_mutex);
    pthread_mutex_unlock(&_is_active_mutex);

    for (it=threads_currently_held.begin(); it!=threads_currently_held.end(); it++){
        if (*it!=current_thread)
            pthread_mutex_unlock(&_spec_threads[*it]._thread_mutex);
    }
}

};

public:

    ///
    /// default constructor
    ///
    loop_speculator():_shared_data((void*)&_null_data){

        _is_active=false;
        _pl_commit=false;
        _has_pre_loop=false;
        _valid_til=-1; //the valid til functionality is not used
        _null_data=-1;
        pthread_mutex_init (&_is_active_mutex, NULL);
        pthread_mutex_init (&_pl_mutex, NULL);
        pthread_mutex_init (&_spec_threads_mutex, NULL);
        pthread_mutex_init (&_valid_til_mutex, NULL);
        pthread_mutex_init (&_global_access_log_mutex, NULL);

    };

    ///
    /// function providing access to the shared_data as a whole
    ///
    void*& get_shared_data () {

        pthread_mutex_lock(&_is_active_mutex);

        bool manages_pre_loop=_has_pre_loop;
        if (!_is_active){

            pthread_mutex_unlock(&_is_active_mutex);

            return (void*)&_null_data; //the object is inactive

        }

        pthread_mutex_unlock(&_is_active_mutex);

        pthread_mutex_lock(&_spec_threads_mutex);

        if (_thread_index.empty() || _thread_index.find(pthread_self())!=_thread_index.end()){

            pthread_mutex_unlock(&_spec_threads_mutex);

            return _shared_data; //valid thread or thread in deferred cancelation

        }
        pthread_mutex_unlock(&_spec_threads_mutex);

        if (manages_pre_loop){

            pthread_mutex_lock(&_pl_mutex);

            if (pthread_self()==_pl){

                pthread_mutex_unlock(&_pl_mutex);

                return _shared_data;

            }

            else {

                pthread_mutex_unlock(&_pl_mutex);

                return (void*)&_null_data; //invalid caller

            }

        }
    }

```

```

        return _shared_data; //unmanaged pre-loop or thread in deferred cancelation
    };

    ///
    /// required function that signals that a given thread has ended
    /// it's execution
    ///
    void commit () {

        pthread_mutex_lock(&_is_active_mutex);

        if (_is_active){

            bool manages_pre_loop=_has_pre_loop;
            pthread_mutex_unlock(&_is_active_mutex);

            pthread_mutex_lock(&_spec_threads_mutex);

            if (!_thread_index.empty()){ //thread in deferred cancel.
                pthread_t thread_id=pthread_self();

                if (_thread_index.find(thread_id)!=_thread_index.end()){

                    int pos=_thread_index.find(thread_id)->second;
                    if (pos>=0){ //valid thread
                        pthread_mutex_lock(&_spec_threads[pos]._thread_mutex);
                        _spec_threads[pos]._commit=true;
                        pthread_mutex_unlock(&_spec_threads[pos]._thread_mutex);
                    }
                }
                else if (manages_pre_loop){
                    if (thread_id==_pl){
                        pthread_mutex_lock(&_pl_mutex);
                        _pl_commit=true;
                        pthread_mutex_unlock(&_pl_mutex);
                    }
                }
                else { //unmanaged pre-loop
                    pthread_mutex_lock(&_pl_mutex);
                    _pl_commit=true;
                    pthread_mutex_unlock(&_pl_mutex);
                }
            }

            pthread_mutex_unlock(&_spec_threads_mutex);
        }
        else{
            pthread_mutex_unlock(&_is_active_mutex);
        }
    };

    ///
    /// function to be called from the speculative threads
    /// and the concurrent pre-loop section, in order to read the shared_data
    /// while keeping the expected sequential data consistency
    ///
    template <class T>
    T read_data(T*& data_to_be_read){

        pthread_mutex_lock(&_is_active_mutex);

        if (_is_active){

            bool manages_pre_loop=_has_pre_loop;
            pthread_mutex_unlock(&_is_active_mutex);

            int current_thread=-2;

            pthread_mutex_lock(&_spec_threads_mutex);

            pthread_t thread_id=pthread_self();

            if (_thread_index.empty()){ //thread in deferred cancel.
                current_thread=-3;
            }
            else if (_thread_index.find(thread_id)!=_thread_index.end()){
                current_thread=_thread_index.find(thread_id)->second;
                if (current_thread==1) //thread in deferred cancel.
                    current_thread=-3;
            }
            else {
                if (manages_pre_loop){
                    pthread_mutex_lock(&_pl_mutex);
                }
            }
        }
    };

```

```

        if (thread_id==_pl){
            current_thread=-1;
        }
        else {
            current_thread=-3; //invalid caller
        }
        pthread_mutex_unlock(&_pl_mutex);
    }
    else {
        current_thread=-1;
    }
}
if (current_thread>=-1){//is a valid read.

    T retval;

    pthread_mutex_lock(&_valid_til_mutex);

    if (_valid_til!=-1 && _valid_til<current_thread){
        retval=*(T*)data_to_be_read;
        pthread_mutex_unlock(&_valid_til_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);
        return retval; //thread that is no longer valid.
    }
    pthread_mutex_unlock(&_valid_til_mutex);

    pthread_mutex_lock(&_global_access_log_mutex);
    if (_access_log.find((void*)data_to_be_read)==_access_log.end()){
        //if the data has never been accessed, it's log should be created

        pthread_mutex_t temp;
        pthread_mutex_init (&temp, NULL);
        _data_access_log temp2;

        //since it's the data's first use, it's assumed to be the value in the pre-loop
        //so the data is initialized on that premise
        temp2._size=sizeof(T);
        temp2._writers.push_back((int)-1);
        temp2._previous_values.push_back((void*)*(T*)data_to_be_read);

        if (current_thread!=-1){
            pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
            _spec_threads[current_thread]._read_data.insert((void*)data_to_be_read);
            temp2._readers.insert(current_thread);//the current read is logged
        }

        pthread_mutex_unlock(&_spec_threads_mutex);
        _access_log.insert (pair<void*, _data_access_log>((void*)data_to_be_read, temp2));
        retval=*(T*)data_to_be_read;

        if (current_thread!=-1){
            pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
        }

        pthread_mutex_unlock(&_global_access_log_mutex);
        return retval;
    }

    _data_access_log* ptr=&_access_log.find((void*)data_to_be_read)->second;
    pthread_mutex_lock(&ptr->_data_mutex);
    pthread_mutex_unlock(&_global_access_log_mutex);

    retval= (T) ptr->_get_previous_value(current_thread);
    vector<int> threads_to_cancel= ptr->_cancel_higher_writers(current_thread);

    //if the log exists, then anti-dependence violation is checked (WAR)

    if (current_thread!=-1){
        pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
    }

    if (!threads_to_cancel.empty()){//an anti-dependence violation has occurred
        memcpy((void*)data_to_be_read, (void*)&retval, sizeof(T)); //the previous value is restored.
        _reset_spec_threads(current_thread, (void*)data_to_be_read, threads_to_cancel);
    }
    else{
        retval=*(T*) data_to_be_read;//no anti-dependence violation has occurred, the data can be read on i
    }
    pthread_mutex_unlock(&_spec_threads_mutex);

    //now that the reading is done, it has to be logged
    if (current_thread!=-1){
        ptr->_readers.insert(current_thread);
        _spec_threads[current_thread]._read_data.insert((void*)data_to_be_read);
    }
}

```

```

        pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
    }
    pthread_mutex_unlock(&ptr->_data_mutex);
    return retval;
}
pthread_mutex_unlock(&_spec_threads_mutex);
}
else {
    pthread_mutex_unlock(&_is_active_mutex);
}
return *(T*) data_to_be_read; //invalid caller, inactive object or thread in deferred cancel.
//it's calling data is returned instead of NULL, to prevent segmentation faults.
};

//!
//! function to be called from the speculative threads
//! and concurrent pre-loop section, in order to write the shared_data
//! while keeping the expected sequential data consistency
//!
template <class T>
int write_data(T*& data_to_be_written_upon, T* data_to_write){

    pthread_mutex_lock(&_is_active_mutex);
    if (_is_active){
        bool manages_pre_loop=_has_pre_loop;
        pthread_mutex_unlock(&_is_active_mutex);

        int current_thread=-2;

        pthread_mutex_lock(&_spec_threads_mutex);

        pthread_t thread_id=pthread_self();

        if (_thread_index.empty()){//thread in deferred cancel.
            current_thread=-3;
        }
        else if (_thread_index.find(thread_id)!=_thread_index.end()){
            current_thread=_thread_index.find(thread_id)->second;
            if (current_thread==-1)//thread in deferred cancel.
                current_thread=-3;
        }
        else {
            if (manages_pre_loop){
                pthread_mutex_lock(&_pl_mutex);
                if (thread_id==_pl){
                    current_thread=-1;
                }
                else {
                    current_thread=-3; //invalid caller
                }
                pthread_mutex_unlock(&_pl_mutex);
            }
            else {
                current_thread=-1;
            }
        }
        if (current_thread>=-1){
            pthread_mutex_lock(&_valid_til_mutex);
            if (_valid_til!=-1 && _valid_til<current_thread){
                pthread_mutex_unlock(&_valid_til_mutex);
                pthread_mutex_unlock(&_spec_threads_mutex);
                return -1;
            }
            pthread_mutex_unlock(&_valid_til_mutex);

            pthread_mutex_lock(&_global_access_log_mutex);
            if (_access_log.find((void*)data_to_be_written_upon)==_access_log.end()){
                //if data has no log entry, then it's entry should be created.

                pthread_mutex_t temp;
                pthread_mutex_init (&temp, NULL);
                _data_access_log temp2;

                //since it's the data's first use, it's assumed to be the value in the pre-loop
                //so the data is initialized on that premise
                temp2._size=sizeof(T);
                temp2._writers.push_back((int)-1);
                temp2._previous_values.push_back((void*)*(T*)data_to_be_written_upon);

                if (current_thread!=-1){
                    pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
                    _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
                    temp2._writers.push_back(current_thread);
                    temp2._previous_values.push_back((void*)*(T*)data_to_be_written_upon);
                }
            }
        }
    }
}

```

```

        else{
            temp2._writers.push_back(-1);
            temp2._previous_values.push_back((void*)data_to_write);
        }
        pthread_mutex_unlock(&_spec_threads_mutex);

        _access_log.insert (pair<void*, _data_access_log>((void*)data_to_be_written_upon, temp2));

        //and finally the writting is done
        memcpy ((void*)data_to_be_written_upon, (void*)&data_to_write, sizeof(T));

        if (current_thread!=-1){
            pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
        }
        pthread_mutex_unlock(&_global_access_log_mutex);
        return 0;
    }

    //if the log exists, then possible true dependence and output dependence violations are checked (RAW and WAW)

    _data_access_log* ptr=&_access_log.find((void*)data_to_be_written_upon)->second;
    pthread_mutex_lock(&ptr->_data_mutex);
    pthread_mutex_unlock(&_global_access_log_mutex);

    if (current_thread!=-1){
        pthread_mutex_lock(&_spec_threads[current_thread]._thread_mutex);
    }

    T value;
    vector<int> writers_to_cancel;

    bool no_output_violation=false;
    if ((void*)*(T*)data_to_be_written_upon==(void*)data_to_write){
        //Since restarting threads has such a significant cost, it is better to check if it is really necessary to
        //So, we will evaluate false positives in output dependence violations.

        if (!ptr->_output_violations_are_false(data_to_write, current_thread)){

            value= (T) ptr->_get_previous_value(current_thread);
            writers_to_cancel=ptr->_cancel_higher_writers(current_thread);

        }
        else {
            no_output_violation=true;
        }
    }
    else{
        value= (T) ptr->_get_previous_value(current_thread);
        writers_to_cancel=ptr->_cancel_higher_writers(current_thread);
    }

    vector<int> readers_to_cancel=ptr->_cancel_higher_readers(current_thread);

    if (no_output_violation && readers_to_cancel.empty()){
        if (current_thread!=-1){
            _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
            pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
        }
        pthread_mutex_unlock(&ptr->_data_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);
        return 0;
    }

    if (!writers_to_cancel.empty()){
        //an output dependence violation has occurred, it's previous value has to be restored before writting over,
        //those writers have to be restarted
        ptr->_writers.push_back(current_thread);
        ptr->_previous_values.push_back((void*)value);
    }

    /*This branch has been turned into a comment, since the logging for all cases where there are false output
    is done on _output_violations_are_false(..). However, should a unforeseeable case arise where this does not
    (for instance, a particular thread reset pattern), then it should be un-commented.*/

    else if (!no_output_violation){
        //no output dependence violation, it's current value is valid and logged as the previous value of
        //thread write...
        ptr->_writers.push_back(current_thread);
        ptr->_previous_values.push_back((void*)*(T*)data_to_be_written_upon);
    }

    if (!readers_to_cancel.empty()){//a true dependence violation has occurred, those readers should be restart
        writers_to_cancel.insert(writers_to_cancel.end(), readers_to_cancel.begin(), readers_to_cancel.end()
    }

    if (!writers_to_cancel.empty()){
        //the violating threads are restarted

```

```

        _reset_spec_threads(current_thread, (void*) data_to_be_written_upon, writers_to_cancel);
    }
    pthread_mutex_unlock(&_spec_threads_mutex);

    //now some more logging of the write
    if (current_thread!=-1){
        ptr->_writers.push_back(-1);
        ptr->_previous_values.push_back((void*)data_to_write);
    }
    else {
        _spec_threads[current_thread]._written_data.insert((void*)data_to_be_written_upon);
    }

    //and finally the write is made
    memcpy ((void*)data_to_be_written_upon, (void*)&data_to_write, sizeof(T));

    if (current_thread!=-1){
        pthread_mutex_unlock(&_spec_threads[current_thread]._thread_mutex);
    }
    pthread_mutex_unlock(&ptr->_data_mutex);
    return 0;
}
pthread_mutex_unlock(&_spec_threads_mutex);
}
else {
    pthread_mutex_unlock(&_is_active_mutex);
}
return -1; //invalid caller, thread in deferred cancel or inactive object.
};

//!
//! function that permits the pre_loop section to dynamically append
//! a new speculative thread at the end of the array in an on-going
//! speculative execution.
//!
int append(void* (f)(void*), void* consts){
    pthread_mutex_lock(&_is_active_mutex);
    if (!_is_active){
        pthread_mutex_unlock(&_is_active_mutex);
        return -1; //the object is not active
    }
    bool manages_pre_loop=_has_pre_loop;
    pthread_mutex_unlock(&_is_active_mutex);

    bool is_pre_loop=false;

    pthread_mutex_lock(&_spec_threads_mutex);

    if (manages_pre_loop){
        pthread_mutex_lock(&_pl_mutex);
        if (pthread_self()==_pl){
            is_pre_loop=true;
        }
        pthread_mutex_unlock(&_pl_mutex);
    }
    if (!is_pre_loop){
        pthread_t thrd_id=pthread_self();
        if ((!_thread_index.empty()) && (_thread_index.find(thrd_id)==_thread_index.end())){
            if (!manages_pre_loop){
                is_pre_loop=true;
            }
        }
    }
}
if (is_pre_loop){
    pthread_attr_t attr;
    pthread_attr_init (&attr);
    pthread_attr_setschedpolicy(&attr, SCHED_RR);
    int success;

    int current_size=_spec_threads.size();

    pthread_mutex_lock(&_valid_til_mutex);
    if (_valid_til!=-1 && (current_size-1)>=_valid_til){
        pthread_mutex_unlock(&_valid_til_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);
        return -1; //the thread to append does not fit with the valid_til condition
    }
    pthread_mutex_unlock(&_valid_til_mutex);

    _spec_threads.resize(_spec_threads.size()+1);
    pthread_mutex_lock(&_spec_threads[_spec_threads.size()-1]._thread_mutex);
    int i=_spec_threads.size()-1;
    _spec_threads[i]._commit=false;
    _spec_threads[i]._thread_instructions=f;
    _spec_threads[i]._const_args=consts;
}

```

```

        success=pthread_create(&_spec_threads[i]._thread, &attr, f, consts);
        if (success!=0){
            pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);
            _spec_threads.resize(_spec_threads.size()-1);
            pthread_mutex_unlock(&_spec_threads_mutex);
            return -1; //the thread could not be created, thus is not appended
        }
        _thread_index.insert(pair<pthread_t, int>(_spec_threads[_spec_threads.size()-1]._thread, _spec_threads.size()-1));
        pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);
        return 0;
    }
    pthread_mutex_unlock(&_spec_threads_mutex);
    return -1; //invalid caller
};

///
///! function that allows the pre_loop section to dynamically cancel
///! a number of speculative threads at the end of the array in an on-going
///! speculative execution, this is done by indicating until what position
///! the execution remains valid.
///!
///! TO BE NOTED: It can only be called once during all the execution.
///!
int valid_til (int pos){
    if (pos<=-1)
        return -1; //invalid argument

    pthread_mutex_lock(&_is_active_mutex);
    if (!_is_active){
        pthread_mutex_unlock(&_is_active_mutex);
        return -1; //the object is not active
    }
    bool manages_pre_loop=_has_pre_loop;
    pthread_mutex_unlock(&_is_active_mutex);

    pthread_mutex_lock(&_valid_til_mutex);
    if (_valid_til!=-1){
        pthread_mutex_unlock(&_valid_til_mutex);
        return -1; //the valid_til function has already been called
    }
    pthread_mutex_unlock(&_valid_til_mutex);

    bool is_pre_loop=false;

    pthread_mutex_lock(&_spec_threads_mutex);

    if (manages_pre_loop){
        pthread_mutex_lock(&_pl_mutex);
        if (pthread_self()==_pl){
            is_pre_loop=true;
        }
        pthread_mutex_unlock(&_pl_mutex);
    }
    if (!is_pre_loop){
        pthread_t thrd_id=pthread_self();
        if ((!_thread_index.empty()) && (_thread_index.find(thrd_id)==_thread_index.end())){
            if (!manages_pre_loop){
                is_pre_loop=true;
            }
        }
    }

    if (is_pre_loop){
        pthread_mutex_lock(&_valid_til_mutex);
        _valid_til=pos;
        pthread_mutex_unlock(&_valid_til_mutex);

        if ( (static_cast<int>(_spec_threads.size()-1)>pos){
            _cancel_higher_spec_threads (pos);
        }
        pthread_mutex_unlock(&_spec_threads_mutex);
        return 0; //no thread was canceled, the restriction will be
                //imposed through the append function
    }
    pthread_mutex_unlock(&_spec_threads_mutex);
    return -1; //invalid caller
};

///
///! speculate: a complete function that takes orderly the instructions and arguments
///! of the pre-loop section, as well as those of the loop iterations; and starts their
///! speculatively parallel execution, while maintaining the sequential consistency.

```

```

///! This function returns in the &shared_data, the results of the computation.
///!
///! TO BE NOTED:
///! * In this version of the function, the object is in control of the pre-
///! loop section. This means that on invocation, the caller blocks until the pre-
///! loop and all the iterations up to the end or valid_til have committed.
///! * If the pre-loop section does not explicitly commit through the commit function,
///! then the whole results of the speculation will not be communicated to the
///! shared data and the speculation will be canceled.
///! * If any iteration does not commit, then only the results up to it will be taken
///! into account, the remaining iterations will be canceled.
///!
int speculate (void*& shared_data, void* (fpl)(void*), void* const_args_pl, script_vector thread_instructions, vector<void*> cons
    if (thread_instructions.size()!=const_args.size()){
        return -1; //invalid parameters
    }
    pthread_mutex_lock(&_is_active_mutex);
    if (_is_active){
        pthread_mutex_unlock(&_is_active_mutex);
        return -1; //the object is active
    }
    _is_active=true;
    pthread_mutex_unlock(&_is_active_mutex);

    pthread_mutex_lock(&_spec_threads_mutex);
    pthread_mutex_lock(&_global_access_log_mutex);

    vector<int> threads_to_cancel;//a selective reset is performed, waiting for the threads in
    _cancel_spec_threads_and_reset_logs(threads_to_cancel);//deferred cancel to finish, and resetting the
//arrays that the speculation uses.

    pthread_mutex_lock(&_is_active_mutex);
    _has_pre_loop=true;
    pthread_mutex_unlock(&_is_active_mutex);

    pthread_mutex_lock(&_valid_til_mutex);
    _valid_til=-1;
    pthread_mutex_unlock(&_valid_til_mutex);

    pthread_mutex_lock(&_pl_mutex);
    _pl_commit=false;
    pthread_mutex_unlock(&_pl_mutex);

    _shared_data=shared_data;

    pthread_attr_t attr;
    pthread_attr_init (&attr);
    pthread_attr_setschedpolicy(&attr, SCHED_RR);

    //the threads are created
    int success;
    int valid_til_thread=-1;
    _spec_threads.resize(thread_instructions.size());
    pthread_mutex_lock(&_pl_mutex);
    success=pthread_create (&_pl, &attr, fpl, const_args_pl);
    pthread_mutex_unlock(&_pl_mutex);
    if (success!=0){
        _spec_threads.clear();
        pthread_mutex_unlock(&_global_access_log_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);

        pthread_mutex_lock(&_is_active_mutex);
        _is_active=false;
        pthread_mutex_unlock(&_is_active_mutex);

        return -1; //the pre-loop section could not be created
    }

    for (unsigned int i=0; i<thread_instructions.size(); i++){
        pthread_mutex_lock(&_spec_threads[i]._thread_mutex);
    }
    for (unsigned int i=0; i<thread_instructions.size(); i++){

        _spec_threads[i]._thread_instructions=thread_instructions[i];
        _spec_threads[i]._const_args=const_args[i];
        success=pthread_create (&_spec_threads[i]._thread, &attr, _spec_threads[i]._thread_instructions, _spec_threads[i].

        if (success==0){
            _thread_index.insert(pair<pthread_t, int>(_spec_threads[i]._thread, i));
        }

        pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);

        if (success!=0){
            for (unsigned int j=i+1; j<thread_instructions.size(); j++){

```



```

        pthread_mutex_unlock(&_spec_threads[j]._thread_mutex);
    }
    for (unsigned int j=0; j<i; j++){
        threads_to_cancel.push_back(j);
    }
    _cancel_spec_threads_and_reset_logs(threads_to_cancel);

    pthread_mutex_unlock(&_global_access_log_mutex);
    pthread_mutex_unlock(&_spec_threads_mutex);

    pthread_mutex_lock(&_is_active_mutex);
    _is_active=false;
    pthread_mutex_unlock(&_is_active_mutex);
    return -1; // a thread could not be created.
}

}

pthread_mutex_unlock(&_global_access_log_mutex);
pthread_mutex_unlock(&_spec_threads_mutex);

bool commit_made=false;
int joined_but_didnt_commit=0;
do{

    success=pthread_join (_pl, NULL);

    if (success!=0){
        pthread_mutex_lock(&_pl_mutex);
        if (_pl_commit){
            commit_made=true;
        }
        else{
            if (joined_but_didnt_commit<10){
                success=0;
                joined_but_didnt_commit++;
            }
        }

        /*just in case the commit is not logged in a timely fasion, it is tried
        10 times to see if it get's made, if not then the object asumes there was no commitment in the pre-loop.
        Although slightly un-elegant, the validity of this solution has not been disproved by empirical
        tests.*/
        pthread_mutex_unlock(&_pl_mutex);
    }

} while (success==0);

pthread_mutex_lock(&_spec_threads_mutex);

if (!commit_made){
    _cancel_higher_spec_threads(-1);
    pthread_mutex_unlock(&_spec_threads_mutex);

    pthread_mutex_lock(&_is_active_mutex);
    _is_active=false;
    pthread_mutex_unlock(&_is_active_mutex);

    return -1; //the pre-loop did not commit.
}

int i=0;

while (_spec_threads.size()>i){
    pthread_mutex_unlock(&_spec_threads_mutex);

    commit_made=false;
    joined_but_didnt_commit=0;

    do{

        success=pthread_join (_spec_threads[i]._thread, NULL);

        if (success!=0){
            pthread_mutex_lock(&_spec_threads[i]._thread_mutex);
            if (_spec_threads[i]._commit){
                commit_made=true;
            }
            else{
                if (joined_but_didnt_commit<100){
                    success=0;
                    joined_but_didnt_commit++;
                }
            }
        }
    }
}

```

```

        pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);
    }

    /*just in case the commit is not logged in a timely fasion, it is tried
    10 times to see if it get's made, if not then the object asumes there was no commitment in the iteration.
    Although slightly un-elegant, the validity of this solution has not been disproved by empirical
    tests.*/

    } while (success==0);

    if (!commit_made){
        pthread_mutex_lock(&_spec_threads_mutex);
        _cancel_higher_spec_threads(i); //This is perhaps a bit excesive, but for consistency purposes it is manda
        pthread_mutex_unlock(&_spec_threads_mutex);

        pthread_mutex_lock(&_is_active_mutex);
        _is_active=false;
        pthread_mutex_unlock(&_is_active_mutex);

        return -1; //the pre-loop did not commit.
    }

    i++;

    pthread_mutex_lock(&_spec_threads_mutex);

    pthread_mutex_lock(&_valid_til_mutex);
    valid_til_thread=_valid_til;
    pthread_mutex_unlock(&_valid_til_mutex);

    if (valid_til_thread!=-1 && i>valid_til_thread){
        i=_spec_threads.size();
    }

}

pthread_mutex_unlock(&_spec_threads_mutex);

pthread_mutex_lock(&_is_active_mutex);
_is_active=false;
pthread_mutex_unlock(&_is_active_mutex);
return 0;
};

//!
//! speculate: a complete function that takes orderly the instructions and arguments
//! of the loop iterations; and starts their speculatively parallel execution, while
//! maintaining the sequential consistency with an un-managed pre-loop section.
//! This function will return in the &shared_data, the results of the computation, when
//! get_results is called.
//!
//! TO BE NOTED:
//! * In this version of the function, the object is not in control of the pre-
//! loop section. This means that on invocation, the caller only blocks for the creation
//! of the threads, and can resume it's execution as a possible pre-loop, until
//! calling get_results(), when the caller blocks until all valid loop iterations/threads return.
//!
int speculate (void*& shared_data, script_vector thread_instructions, vector <void*> const_args){
    if (thread_instructions.size()!=const_args.size()){
        return -1; //invalid parameters
    }
    pthread_mutex_lock(&_is_active_mutex);
    if (_is_active){
        pthread_mutex_unlock(&_is_active_mutex);
        return -1; //the object is active
    }
    _is_active=true;
    pthread_mutex_unlock(&_is_active_mutex);

    pthread_mutex_lock(&_spec_threads_mutex);
    pthread_mutex_lock(&_global_access_log_mutex);

    vector <int> threads_to_cancel;//a selective reset is performed, waiting for the threads in
    _cancel_spec_threads_and_reset_logs(threads_to_cancel);//deferred cancel to finish, and resetting the
    //arrays that the speculation uses.

    pthread_mutex_lock(&_is_active_mutex);
    _has_pre_loop=false;
    pthread_mutex_unlock(&_is_active_mutex);

    pthread_mutex_lock(&_valid_til_mutex);
    _valid_til=-1;
    pthread_mutex_unlock(&_valid_til_mutex);

```

```

pthread_mutex_lock(&_pl_mutex);
_pl_commit=false;
pthread_mutex_unlock(&_pl_mutex);

_shared_data=shared_data;

pthread_attr_t attr;
pthread_attr_init (&attr);
pthread_attr_setschedpolicy(&attr, SCHED_FIFO);

//the threads are created
int success;
_spec_threads.resize(thread_instructions.size());

for (unsigned int i=0; i<thread_instructions.size(); i++){
    pthread_mutex_lock(&_spec_threads[i]._thread_mutex);
}
for (unsigned int i=0; i<thread_instructions.size(); i++){

    _spec_threads[i]._thread_instructions=thread_instructions[i];
    _spec_threads[i]._const_args=const_args[i];
    success=pthread_create (&_spec_threads[i]._thread, &attr, _spec_threads[i]._thread_instructions, _spec_threads[i].

    if (success==0){
        _thread_index.insert(pair<pthread_t, int>(_spec_threads[i]._thread, i));
    }

    pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);

    if (success!=0){
        for (unsigned int j=i+1; j<thread_instructions.size(); j++){
            pthread_mutex_unlock(&_spec_threads[j]._thread_mutex);
        }
        for (unsigned int j=0; j<i; j++){
            threads_to_cancel.push_back(j);
        }
        _cancel_spec_threads_and_reset_logs(threads_to_cancel);

        pthread_mutex_unlock(&_global_access_log_mutex);
        pthread_mutex_unlock(&_spec_threads_mutex);

        pthread_mutex_lock(&_is_active_mutex);
        _is_active=false;
        pthread_mutex_unlock(&_is_active_mutex);
        return -1;// a thread could not be created.
    }
}
pthread_mutex_unlock(&_global_access_log_mutex);
pthread_mutex_unlock(&_spec_threads_mutex);

return 0;
};

///
/// function to get the results of the speculation once the un-managed pre-
/// loop section has ended it's execution.
///
/// TO BE NOTED:
/// * If the pre-loop section does not explicitly commit through the commit function,
/// then the whole results of the speculation will not be communicated to the
/// shared data and the speculation will be canceled.
/// * If any iteration does not commit, then the object will be in an infinite loop wa-
/// iting for it to commit.
int get_results(){
    pthread_mutex_lock(&_is_active_mutex);
    if (!_is_active || !_has_pre_loop){
        pthread_mutex_unlock(&_is_active_mutex);
        return -1;//the object is inactive or manages it's pre-loop section, thus this function cannot be run.
    }
    pthread_mutex_unlock(&_is_active_mutex);

    bool commit_made=false;
    unsigned int joined_but_didnt_commit=0;
    int success;

    do{

        success=-1;

        if (success!=0){
            pthread_mutex_lock(&_pl_mutex);
            if (_pl_commit){
                commit_made=true;
            }
            else{

```

```

        if (joined_but_didnt_commit<10){
            success=0;
            joined_but_didnt_commit++;
        }
    }
    pthread_mutex_unlock(&_pl_mutex);
}

} while (success==0);

int i=0;

pthread_mutex_lock(&_spec_threads_mutex);

if (!commit_made){
    _cancel_higher_spec_threads(-1);
    pthread_mutex_unlock(&_spec_threads_mutex);

    pthread_mutex_lock(&_is_active_mutex);
    _is_active=false;
    pthread_mutex_unlock(&_is_active_mutex);

    return -1; //the pre-loop did not commit.
}

while (_spec_threads.size()>i){

    pthread_t thread_to_join=_spec_threads[i]._thread;

    pthread_mutex_unlock(&_spec_threads_mutex);

    commit_made=false;
    joined_but_didnt_commit=0;

    do{

        success=pthread_join (thread_to_join, NULL);

        if (success!=0){
            pthread_mutex_lock(&_spec_threads[i]._thread_mutex);
            if (_spec_threads[i]._commit){
                commit_made=true;
            }
            else{
                if (joined_but_didnt_commit<100){
                    success=0;
                    joined_but_didnt_commit++;
                }
            }
            pthread_mutex_unlock(&_spec_threads[i]._thread_mutex);
        }

        /*just in case the commit is not logged in a timely fasion, it is tried
        10 times to see if it get's made, if not then the object asumes there was no commitment in the iteration.
        Although slightly un-elegant, the validity of this solution has not been disproved by empirical
        tests.*/

    } while (success==0);

    if (!commit_made){

        pthread_mutex_lock(&_spec_threads_mutex);
        _cancel_higher_spec_threads(i); //This is perhaps a bit excesive, but for consistency purposes it is manda
        pthread_mutex_unlock(&_spec_threads_mutex);

        pthread_mutex_lock(&_is_active_mutex);
        _is_active=false;
        pthread_mutex_unlock(&_is_active_mutex);

        return -1; //some thread did not commit.
    }

    i++;

    pthread_mutex_lock(&_spec_threads_mutex);

    pthread_mutex_lock(&_valid_til_mutex);
    int valid_til_thread=valid_til;
    pthread_mutex_unlock(&_valid_til_mutex);

    if (valid_til_thread!=-1 && i>valid_til_thread){
        i=_spec_threads.size();
    }
}

```

```
    }

    pthread_mutex_unlock(&_spec_threads_mutex);

    pthread_mutex_lock(&_is_active_mutex);
    _is_active=false;
    pthread_mutex_unlock(&_is_active_mutex);
    return 0;
};

};
```